

Stress-Testing the Persistent Asset Charge

Does the Architecture Survive Contact with Reality?

Drew Griffiths

This is the companion stress-testing pass to the architecture paper. Its purpose is not to defend the model but to try to break it, and to report honestly what breaks. Each test is classified on a four-way scale: fatal (the model cannot work as designed), serious but addressable (a real problem with a known fix that must be adopted), manageable (a cost or limit to be accepted and disclosed), or unresolved (genuinely open, needing work this pass cannot complete). The two tests most likely to be fatal — revenue sufficiency and human-rights durability — are taken first, because if either failed, the rest of the pass would be beside the point.

Neither returned fatal. On the evidence available to this review, no fatal defect has been identified, which is a more modest statement than that none exists. But the revenue test now returns an adverse result of a different kind, and the distinction matters: this pass adopted full replacement of Capital Gains Tax and Inheritance Tax as a deliberately severe condition, and on current receipts the model does not clear it. That is a failure against a ceiling this pass set for itself, not a disqualification on revenue, and Section 1 keeps the two apart. The model's viability remains conditional on specific design safeguards identified below and on external work — a government costing and counsel's opinion — that this pass cannot complete.

Version 2, 29 August 2026 (first published 21 July 2026)

Anyone asked to support this proposal will ask two questions before any other: will it raise the money, and will it survive legal challenge. If the answer to either is no, nothing else matters. This pass therefore leads with those two, states its findings in the harshest defensible terms, and only then proceeds to the secondary tests. Throughout, the burden of proof is placed on the model: a test is not passed because the model can be defended, but because a serious critic would concede the point.

A word on method. This pass uses the best publicly available UK data — HM Revenue and Customs receipts, the Office for Budget Responsibility's forecasts, the Wealth Tax Commission's base estimates, and the European Court of Human Rights case-law on taxation and property. It does not have access to the confidential taxpayer-level data on which a government costing would rest. Where a finding depends on data only government holds, that is stated, and the finding is marked accordingly. No number in this pass should be treated as a costing; the numbers are order-of-magnitude tests of whether the model is in the right region, not forecasts.

Confidence and limits of this pass, stated once and plainly. Four things a reader relying on this document must know up front. First, the revenue figures rest on soft data: the count of people above £10m is reasonably robust, but the aggregate wealth they hold is under-captured at the very top by the survey sources it derives from, and the base figure used here is a consistency check against advocacy estimates, not an independent measurement — a circularity disclosed in Section 1.2. Second, the entity-layer contribution to revenue is directionally positive but unquantified in this pass; the public data does not permit sizing it, and no figure for it should be relied upon. Third, the legal findings are conditional on safeguards the architecture paper did not originally contain: the accurate statement is that reciprocal purchase is capable of being made compliant, not that it was compliant as first drafted. Addressable-in-principle is not addressed-in-fact. Fourth, the replacement condition tested at Section 1 is one this pass imposed; the architecture paper presents replacement as a possibility rather than an objective, and failing that condition is not the same as the model raising too little.

1. Test 1 — Revenue: does it clear the replacement ceiling?

Classification: does not clear, on current receipts. At the tested rates of 1.5 to 2 per cent the model does not clear the full-replacement ceiling this pass adopts. Clearing it would require roughly 2.8 to 3.9 per cent on the personal base, a region in which the coupled proportionality problem at Test 3 becomes materially worse. **Failing this ceiling is not disqualification on revenue**, because full replacement is a condition this pass imposed and not one the architecture asserts.

1.1 The ceiling, and why this pass adopts it

The architecture paper does not claim to replace Capital Gains Tax and Inheritance Tax. It says the model is not a replacement for them "dressed up as something grander (though it could replace them)", and reserves revenue sufficiency as an empirical question for this pass.

This pass nonetheless adopts full replacement of CGT and IHT as its revenue test, and the reason is adversarial rather than definitional. The burden of proof here is on the model, and a condition a serious critic would concede is a stronger test than one the model was designed to meet. Full replacement is therefore a **ceiling this pass sets**, not a requirement the architecture must satisfy. What follows should be read accordingly: clearing it would be strong evidence, and failing it withdraws that evidence without establishing that the model raises too little.

Which series, and why it matters. "CGT revenue" names two different quantities. *Receipts* are cash collected in a year; *liabilities* are tax assessed on gains realised in a year, collected later. For 2024-25 the

two differ by £10.5bn on the same tax: liabilities of £24.2bn against receipts of £13.7bn. This pass uses **cash receipts throughout**, because the replacement question is what the Exchequer would forgo in cash terms, and because IHT is published on that basis. The cost of that choice is that a receipts series carries payment-timing volatility, which the figures below illustrate.

- **CGT:** £14.5bn in 2023-24 and £13.7bn in 2024-25, rising sharply to **£24.3bn** in 2025-26. The OBR attributes the surge to disposals made in 2024-25 to benefit from lower rates ahead of anticipated policy changes, whose CGT fell due mainly in January 2026, and forecasts further growth to £34.9bn by 2030-31 driven principally by higher equity prices.
- **IHT:** £7.5bn in 2023-24, £8.2bn in 2024-25, and a record £8.5bn in 2025-26, forecast to keep rising as thresholds stay frozen.

All four figures are taken from HMRC's own statistics table, Table 2, released 21 August 2026, on a cash basis, with the 2025-26 year aligned to the HMRC annual report and accounts published on 9 July 2026.

The ceiling is therefore **£32.8bn** on the 2025-26 pair, being £24.310bn of CGT and £8.467bn of IHT, against the £21–23bn this pass used when first written, and it is projected to rise further. A replacement matching the earlier band would already be some £10bn short.

A correction to the earlier reading. The first version of this pass treated the 2024-25 fall in CGT receipts as evidence of realisation-based lock-in. That reading is withdrawn. The fall was a payment-timing effect of a rate change announced in October 2024, and liabilities for the same year rose 89 per cent to £24.2bn from 584,000 taxpayers, as holders made disposals to benefit from lower rates ahead of anticipated policy changes. That is evidence of *timing* behaviour, which is the architecture's second avoidance channel, not its first. Lock-in remains well supported in the wider literature; this particular observation cannot carry it.

The two series corroborate that reading independently. Liabilities assessed on 2024-25 gains were £24.2bn; cash receipts in 2025-26 were £24.310bn. The near equality is what the timing account predicts, since most of that liability fell due in January 2026. It is strong support for the explanation rather than proof that every pound in one series maps to the other, and it is recorded as support.

On classification. This result carries no entry on the pass's four-way severity scale, and that absence is deliberate. Those four classes answer what kind of defect the model has encountered. This answers what happened when the model was tested against a deliberately stronger condition. The axes are orthogonal and should stay so: **defect classification** is fatal, serious but addressable, manageable or unresolved; **adversarial-condition outcome** is clears, does not clear, or unresolved.

1.2 The base

The base rests on two anchors of very different strength, and honesty requires separating them.

The first anchor is the population count: independent estimates converge on roughly 20,000 to 32,000 individuals with personal net wealth above £10m. This figure is reasonably robust — several sources reach the same region by different routes — and can be leaned on.

The second anchor is the aggregate wealth that population holds, and it is much softer. The Wealth Tax Commission and subsequent analyses imply that the excess of this group's wealth above the £10m threshold aggregates to approximately £1.0–1.2 trillion. But this figure is derived by back-solving from advocacy estimates — that a charge of around 2 per cent on the excess would raise about £24bn implies an excess base of about £1.2 trillion — and those estimates rest partly on survey data known to under-capture wealth at the very top. Using a figure in that region to then test this model is therefore a consistency check against the advocacy numbers, not an independent measurement. The circularity is disclosed rather than hidden: the true base could be materially larger (top-end under-capture would

mean more wealth than the surveys show, helping revenue) or the effective base materially smaller (behavioural response and valuation disputes could shrink it, hurting revenue). No weight should be placed on the precise trillion figure; the honest claim is only that the base is of an order that makes the model worth testing, not that it is known.

The proposed model is a cliff, not an allowance: once the threshold is crossed, the full averaged value is charged, not merely the excess. This adds back the first £10m for each liable person — roughly £200–320bn across the population, this increment resting on the more robust population count — giving a personal cliff base of the order of £1.2–1.5 trillion, subject to the same disclosed softness in the underlying wealth figure, and before the entity layer is counted at all.

1.3 The arithmetic against the ceiling

On the mid-estimate personal cliff base of £1.3 trillion, with three behavioural haircuts of 10 per cent (optimistic), 20 per cent (central, the Wealth Tax Commission's working assumption) and 35 per cent (pessimistic, reflecting a first-of-its-kind annual charge with no exemptions):

rate	gross	net at 10%	net at 20%	net at 35%
2.0%	£26.0bn	£23.4bn	£20.8bn	£16.9bn
2.5%	£32.5bn	£29.2bn	£26.0bn	£21.1bn
3.0%	£39.0bn	£35.1bn	£31.2bn	£25.4bn
3.5%	£45.5bn	£41.0bn	£36.4bn	£29.6bn

Against the ceiling of £32.8bn, the tested 2 per cent fails under every haircut, and 3 per cent clears it only under the optimistic one, netting £25.4bn to £35.1bn. The rate required to clear it is **2.8 per cent under the optimistic haircut, 3.2 per cent under the central one, and 3.9 per cent under the pessimistic one.**

1.4 The entity layer is the model's revenue advantage — and its accounting hazard

The figures above count only personally-held wealth. The model's distinctive feature is that it also charges companies, trusts, and other liable persons independently. That entity layer is additional base not captured by conventional personal wealth taxes, and it is the reason the model can, in principle, out-raise them. This is a genuine structural advantage.

But it cannot simply be added to the personal figure, and honesty requires saying so plainly. Where a shareholder's shares and the company's underlying assets are both charged, the same economic value appears at both levels. How much net revenue the entity layer adds therefore depends entirely on the interaction rule — the very question the architecture paper left, correctly, to Parliament:

- under **full stacking**, the entity layer adds a substantial increment, because the same value is charged at both levels — but at the cost of taxing that value twice in full, which is where the human-rights proportionality question (Test 2) bites hardest;
- under a **credit or offset**, related value is charged once at the higher level, the entity layer adds little, and revenue rests mainly on the personal base.

A claim in an earlier draft of this pass — that full stacking "comfortably clears the hurdle at 1 per cent" — is withdrawn as unestablished. Sizing the entity layer requires knowing how much of the £10m-plus personal wealth is held through UK-resident companies and trusts rather than directly, at a granularity the public data does not provide. The honest statement is therefore that the entity layer's net

contribution is **directionally positive but unquantified in this pass**, and no revenue figure should rest on it.

That refusal now costs more than it did. At 1.5 to 2 per cent this pass could leave the entity layer unsized without consequence, because the personal base was carrying the test. It no longer is. An unquantified component would now be required to rescue a quantified principal test, and this pass will not permit that: the position is unchanged and the figure stays out.

This is the single most important interaction in the whole model between the revenue question and the legal question: the rate the model needs and the legal risk it runs move in opposite directions as the interaction rule changes. Full stacking raises more and raises proportionality risk; crediting lowers legal risk and lowers the entity layer's revenue contribution. This trade-off should be stated openly to anyone considering the model, because it cannot be designed away — and, as Section 3 notes, it is not the only place where a revenue dial is also a legal-risk dial.

1.5 Verdict on revenue

Three findings, stated separately because they are separate.

1. **The model produces substantial recurring revenue.** On the personal base alone, at 1.5 to 2 per cent, it is of the order of £13bn to £23bn net. That is a material instrument by any measure.
2. **It does not clear the full-replacement ceiling on current receipts.** At 2 per cent it falls short of £32.8bn under every behavioural assumption tested, at 3 per cent it clears only under the optimistic haircut, and the ceiling is rising.
3. **Clearing the ceiling would require a rate the model's other tests make harder to defend.** Test 3 identifies the rate as simultaneously a revenue dial and a legal-risk dial. A charge of 3.2 per cent on the full averaged value, not the excess, takes most of the real return of a holder earning 3 per cent real, which strengthens the individual-and-excessive-burden argument precisely where the revenue calculation now points.

Three cautions from the original pass remain binding. The base estimate is the Wealth Tax Commission's, built partly on survey data known to under-capture the very top, so the true base could be larger (helping) or the behavioural response larger (hurting). Only government, with taxpayer-level data, can convert an order-of-magnitude test into a costing. And the entity layer, whatever its true size, cannot be counted until it is measured.

What follows from those three findings is a question this pass cannot settle and should not disguise: whether any rate that clears the ceiling survives the coupled tests. That is now this pass's principal open question on the replacement ceiling, and it has replaced the earlier one.

2. Test 2 — Human-rights durability: does the charge, and reciprocal purchase, survive A1P1 and Article 6?

Classification: serious but addressable, with one component (reciprocal purchase) that is conditionally survivable only if specific safeguards from the case-law are built in. Without those safeguards it is fatal to that mechanism; with them it is defensible.

(Sections 2.1 to 2.5 are unchanged from version 1. The architecture paper has since adopted both conditions at its Section 4.3, which is recorded in the correction note at the end of this document; the finding itself is unaltered.)

2.1 The favourable baseline

The starting position is more favourable than instinct suggests. Under Article 1 of Protocol 1 (A1P1, protection of property), the European Court of Human Rights affords the legislature a wide margin of appreciation in tax matters. Few cases succeed under Article 1 in the tax field; where they do, there is an element of clear arbitrariness, and they have almost all arisen in the enforcement field rather than the imposition of a charge. There is no general principle against even retrospective taxation. An annual wealth charge, as such, is well within the margin the Court allows. So the existence of the charge is not the problem.

2.2 The operative constraint: individual and excessive burden

The binding A1P1 constraint is the requirement of a fair balance between the general interest and the individual's rights, breached where a measure imposes an "individual and excessive burden." This is where two features of the model — neither of them the charge itself — attract risk:

- **The cliff.** Charging the full value once the threshold is crossed, rather than the excess, produces a discontinuity that could, at the margin, be argued to impose an excessive individual burden on someone just over the line. This is manageable (see Test 3) but should be understood as the model's most exposed ordinary feature.
- **The illiquid holder.** A person whose averaged wealth exceeds the threshold but who lacks the cash to pay — the private-company or farm owner — is the classic "excessive burden" claimant if forced to sell to pay an annual charge. The architecture's answer (liability determined normally, but payment deferrable with interest) is the correct one and is close to essential for A1P1 compliance; it should be treated not as an optional kindness but as a durability requirement.

2.3 Reciprocal purchase and the Hentrich problem

The sharpest legal finding of this pass concerns reciprocal purchase, and it comes from *Hentrich v France* (1994). France operated a revenue right of pre-emption: where the authorities considered a declared sale price too low, they could buy the property at the declared price. The Court found a violation of A1P1 — and a linked violation of Article 6 (fair trial) — but the reasons are what matter, because they map precisely onto the model's design choices.

The violation was not that state pre-emption at a declared value is inherently incompatible with property rights. It was that:

1. the applicant was a "selected victim" of a discretionary power applied to her individually and arbitrarily; and
2. she had no effective opportunity to challenge the measure — the procedure did not let her contest the authorities' view or present her case on equal terms.

The Court expressly said the burden could have been rendered legitimate only if she had had the possibility, which was refused her, of effectively challenging the measure.

This is the decisive conditional for the model. Reciprocal purchase is Hentrich-compliant only if it is built to avoid exactly the two defects the Court identified:

- **It must not be arbitrary or selective.** The power must operate under clear, general, published rules applying uniformly to all who make a substitute-value election — not as a discretionary power the state aims at chosen individuals. The architecture's design helps here: the power is triggered by the taxpayer's own election of a low value, not by the state selecting a victim. That is a materially stronger position than Hentrich, where the state initiated. But the rules governing when and how the

state may exercise the resulting option must be general and transparent, or the Hentrich defect reappears.

- **It must afford an effective challenge and equality of arms.** The taxpayer must have a genuine, timely, adequately-resourced route to contest valuation and process, on terms that do not place them at a substantial disadvantage against the state. This is why the architecture's separation of the ordinary challenge (no purchase exposure) from the substitute-value election (purchase exposure) matters: the ordinary challenge route is what supplies the effective possibility of challenging that Hentrich required and France lacked.

There is also an Article 6 length-of-proceedings point in Hentrich (the French proceedings ran over seven years). The architecture's "pay throughout, correct forward with interest" rule must be paired with a genuinely expeditious challenge process, or a long dispute during which the taxpayer pays on a contested value could itself found an Article 6 complaint.

2.4 The second, independent condition: Article 6 and the timing of a contested charge

There is a further requirement in Hentrich that is distinct from non-selectivity and deserves to be treated as a separate condition rather than a sub-clause of it. The Hentrich proceedings ran over seven years, and the Court found a violation of Article 6 on the length alone. This matters acutely for the present model because its design deliberately runs the charge during a dispute: under the "pay throughout, correct forward with interest" rule, a taxpayer contesting a valuation pays on the state's figure for as long as the dispute lasts. A perfectly non-selective, rule-based reciprocal purchase could still generate an Article 6 violation purely through delay, if that delay is long and the taxpayer is paying on a contested value throughout.

So there are two independent conditions on the mechanism, not one:

1. **Non-selectivity and effective challenge** (the A1P1 condition): rule-based, uniform, with a genuine route to contest.
2. **Expeditious resolution** (the Article 6 condition): the challenge and any purchase process must be genuinely fast, precisely because the charge accrues on a contested value while the dispute runs. A slow process is not merely inconvenient; on the authority of Hentrich it is independently unlawful.

A supporter should see both, because a design that satisfied the first and failed the second would still be struck down.

2.5 Honest verdict on durability

The charge survives. Reciprocal purchase is **capable of being made compliant** — but the accurate statement is precisely that, not that it was compliant as the architecture originally stood. The architecture paper did not contain the anti-selectivity rules or the expeditious-resolution guarantee; it described the mechanism but not the safeguards that make it lawful. So the correct finding is: the mechanism can be drafted to survive A1P1 and Article 6, and because its trigger is the taxpayer's own election rather than state selection it starts from a stronger position than Hentrich — but it survives only if both conditions above are written into the charging legislation, and a version lacking either would fail. These are not refinements; they are the difference between a compliant mechanism and a struck-down one, and they are exactly the points a competent tax silk will raise first. This must be disclosed to anyone supporting the model.

3. Test 3 — The cliff and the excessive-burden margin

Classification: manageable, provided quarterly averaging and payment deferral are retained and disclosed.

The entry-threshold (cliff) design is a deliberate choice, defended in the architecture paper as a trade-off for a broader base. Under stress it holds, but only with its accompanying safeguards treated as load-bearing:

- quarterly averaging prevents a single day's valuation from determining liability on a whole holding, which both reduces arbitrariness (helping A1P1) and defeats one-day boundary management;
- payment deferral with interest addresses the liquidity case that would otherwise be the strongest "excessive burden" claim;
- the discontinuity itself remains, and should be disclosed as an accepted cost — a marginal person just over the threshold bears a real jump. This is defensible as a matter of legislative line-drawing (all thresholds draw lines) but it is the model's most criticised ordinary feature and should never be hidden.

A residual honest point: if the excessive-burden risk at the margin proves sharper than expected, the fallback is a taper band just above the threshold, which softens the discontinuity without abandoning the broad base. Adopting a taper would be a legislative choice, not an architectural change, and it is worth naming as the available mitigation.

There is also a connection to Test 1 worth surfacing, because it completes a pattern, and the connection has tightened since this pass was first written. The revenue analysis now finds that clearing the replacement ceiling would require roughly 2.8 to 3.9 per cent on the personal base, against the 1.5 to 2 per cent originally reported. But the rate is not only a revenue dial — it is simultaneously a legal-risk dial, because a higher annual charge strengthens any "individual and excessive burden" argument, most sharply at the cliff edge and for the illiquid holder. So two of the model's key parameters both trade revenue against proportionality risk in the same direction: the interaction rule (Test 1) and the rate (here). Pushing either toward more revenue pushes the model toward more legal exposure. A decision-maker should see this as a single coherent constraint rather than two separate ones: the model's revenue and its human-rights durability are coupled, and cannot both be maximised at once. The higher the required rate, the tighter that coupling becomes.

4. Test 4 — Avoidance migration

Classification: serious but addressable; partly unresolved pending data.

The architecture paper claims resistance to realisation-based avoidance only, and that claim holds: a charge that does not wait for a disposal cannot be escaped by not disposing. But avoidance migrates, and the pass must say where:

- **Beneficial-ownership concealment** becomes the primary avoidance channel. The model's Principle 3 (economic ownership over legal form) is the intended answer, but where ownership is genuinely unattributable — offshore discretionary trusts, layered foreign entities — it has less to bite on. This is the model's real frontier, shared with every wealth tax, and it is only partly addressable by a single jurisdiction. Unresolved to the extent it depends on international beneficial-ownership transparency.

- **Valuation gaming** shifts from "declare low" (disciplined by reciprocal purchase) to disputes over what is included, minority discounts, and control premiums. Manageable through valuation rules, but real.
- **Debt manufacture** is addressed by the bounded-net rule (related-party and artificial liabilities disregarded), which is the correct design and must be robustly drafted.
- **Threshold planning** — fragmenting holdings across connected persons — is addressed by connected-person aggregation, but that aggregation rests on a normative rule (that genuinely independent persons are not aggregated merely by connection) which the architecture paper honestly flagged as needing further specification. Serious but addressable, requiring careful anti-fragmentation drafting.

The honest overall finding: the model removes realisation-based avoidance at its root and forces the remainder into channels that are visible and, mostly, addressable — but the offshore-opacity channel is a genuine residual limit that no domestic design closes, and the model should claim no more.

5. Test 5 — Capital flight and situs

Classification: **manageable** for immovable and UK-situated wealth; **unresolved** for portable financial wealth.

Because liability attaches to assets by situs, UK land, property, and domestic business assets cannot leave, and the charge continues under any owner — so the flight response that undermines conventional wealth taxes is genuinely blunted for the immovable majority of the base. This is a real structural advantage and survives stress.

The residual is portable financial wealth — listed securities and cash — which a departing owner can move. A charge on realised value leaving the jurisdiction is a partial answer; full capture is beyond any single state. This is a genuine limit, not a fatal one, because the immovable base is large and anchored; but the model should not claim to solve flight for portable wealth, and this pass does not.

6. Test 6 — Administrative feasibility

Classification: **serious but addressable; a multi-year build with real cost.**

A continuously-maintained register of beneficial ownership and standing values across 20,000–32,000 liable persons, plus entities, plus those near the threshold who must file, is a substantial institutional undertaking. The Wealth Tax Commission estimated setup costs of the order of £580m for a conventional annual wealth tax; this model, with its register, valuation function, reciprocal-purchase capability, and disposal machinery, is at least as demanding and probably more so.

Two specific requirements are load-bearing rather than incidental:

- **A credible, funded state purchase-and-disposal capacity.** Reciprocal purchase disciplines valuation only if the state can and would exercise it. Without a standing, funded acquisition-and-disposal function, the discipline is decorative — and, as the architecture paper noted, acquisition of operating businesses can itself destroy value, so the capability must be asset-class-differentiated. This is a real institutional cost that must be budgeted, not assumed.

- **Valuation capacity at scale.** Setting and defending standing values for illiquid assets — private companies, land, art — is the function that has historically overwhelmed wealth-tax administrations. The burden-shift and reciprocal-purchase design is the model's answer, and it is a good one, but it does not eliminate the need for a serious valuation capability; it changes its shape.

Feasible, then, but only as a multi-year build with a real budget, and no honest presentation should imply otherwise.

7. Test 7 — Cumulative taxation and fairness

Classification: manageable; a designed feature to be defended, not denied.

The model deliberately charges the same underlying value at more than one level (company and shareholder). The correct defence, confirmed by this pass, is not to deny cumulative taxation but to locate it: the charges fall on different liable persons for a distinct legislative purpose, exactly as corporation tax and dividend tax already do on the same corporate profits. That precedent is directly on point and is a strong answer. The fairness question — whether the resulting aggregate burden is justified — is a political one the model does not pretend to settle, and the interaction-rule choice (Test 1) is where it is actually decided.

8. What would make a supporter safe — and what would not

Given the stakes in putting one's name to this, the pass ends with a direct statement of what a supporter can and cannot safely claim.

The one sentence to carry. If a supporter remembers nothing else from this section:

The architecture never claimed to replace CGT and IHT in full; this pass deliberately set full replacement as a harder ceiling, and on current receipts the tested rate does not clear it.

A supporter can safely say:

- that the model is a genuine, original answer to realisation-based avoidance and to the legal-form escape, not a re-packaged wealth tax;
- that it raises **substantial recurring revenue** — of the order of £13bn to £23bn net on the personal base at 1.5 to 2 per cent, with the entity layer as directionally positive but unsized additional capacity;
- that it **does not, on current receipts, establish full replacement of CGT and IHT** — a ceiling this pass set for itself rather than one the architecture claims, and one that has risen materially since this pass was first written;
- that it is not disqualified on human-rights grounds — the charge sits within the wide margin the ECtHR allows, and reciprocal purchase is capable of being made compliant if built with the Hentrich safeguards, which the architecture paper now specifies;
- that its limitations are disclosed rather than hidden, which is itself a mark of seriousness.

A supporter must not claim:

- **that failing the replacement ceiling disqualifies the model** — full replacement is a condition this pass imposed, and the architecture reserved revenue sufficiency rather than asserting it;

- **that raising the rate to clear the ceiling is a free adjustment** — the rate is also a legal-risk dial, and the two move together;
- that it raises a specific sum — only government can cost it, the base rests on soft top-end data, and the honest statement is "in the right region, rate-dependent, subject to behavioural response";
- that the entity layer raises a quantified amount — its contribution is directionally positive but unsized in this pass, and the "clears at 1 per cent under stacking" claim is withdrawn as unestablished;
- that it is airtight against avoidance — offshore opacity remains a real residual limit;
- that reciprocal purchase is safe without the safeguards — it is capable of being made safe, and only once the non-selectivity and expeditious-resolution conditions are written into the charging legislation; a version without them would fail on Hentrich;
- that the cliff has no cost — it has a real discontinuity, defensible but not costless;
- that revenue and legal safety can both be maximised — the interaction rule and the rate each trade one against the other;
- that it can be stood up quickly or cheaply — it is a multi-year institutional build.

9. Questions now answered — and questions still open

It is worth separating, at the end, what this pass has settled in principle from what still requires evidence it could not obtain. The distinction matters because a supporter needs to know which objections they can already meet and which remain genuinely live.

Answered in principle by this pass:

- the model's constitutional identity — a persistent charge on the holding of concentrated assets by any liable person, distinct from event-driven taxes;
- that the personal base at 1.5 to 2 per cent produces substantial recurring revenue, of the order of £13bn to £23bn net;
- that the charge itself falls within the ECtHR's wide margin of appreciation in tax;
- that reciprocal purchase has a route to compliance, via the taxpayer-election trigger and the Hentrich safeguards;
- that the principal avoidance channel it removes (realisation-based) is genuinely closed, and where the residue migrates.

Answered adversely by this pass:

- that the model does **not**, at the rates tested and on current receipts, clear the full-replacement ceiling this pass adopted; and that clearing it would require a rate in the region where the coupled proportionality problem worsens.

Still requiring external evidence this pass could not supply:

- a government revenue costing on taxpayer-level data, including the size of the entity layer, which the public data cannot size;
- counsel's opinion on reciprocal purchase and on the cliff's excessive-burden exposure, which a case-law reading can frame but not settle;
- an administrative costing of the register, valuation function, and state purchase-and-disposal capacity;

- behavioural-response modelling specific to this design rather than borrowed from lower-threshold historic wealth taxes;
- the transition path from CGT and IHT, which this pass has not examined at all;
- whether any rate that clears the replacement ceiling survives the coupled tests, which is now the principal open question on that ceiling.

10. Overall finding

The architecture survives the stress-testing pass conditionally, and the word conditionally is doing real work. No test returned fatal. On the evidence available to this review, no fatal defect has been identified — which is a different and more modest statement than that none exists, because several of the findings rest on data and law this pass could frame but not finally settle.

Two tests returned "serious but addressable" on the load-bearing question of durability and on avoidance migration, each with a specific, nameable condition: reciprocal purchase built with both the non-selectivity and the expeditious-resolution safeguards, and careful anti-fragmentation drafting. The remainder are manageable costs or genuine but bounded limits, with offshore opacity the one materially unresolved frontier, shared with every wealth tax and not closable by domestic design alone.

The revenue test returns an outcome of a different kind, and this pass keeps it distinct. **The model does not clear the full-replacement ceiling on current receipts.** That ceiling was chosen by this pass as a deliberately severe adversarial condition, not asserted by the architecture, which presents replacement as a possibility and reserves sufficiency as an open question. Failing it withdraws a strong piece of evidence; it does not establish that the model raises too little, and on the personal base at the tested rates the model raises of the order of £13bn to £23bn net. What it does establish is that the distance between the model's demonstrated revenue capacity and full replacement is larger than this pass originally reported, and that closing it runs through parameters the model's other tests make harder to defend.

The honest one-sentence verdict, suitable for someone deciding whether to put their name to it: on the evidence available to this review, this is a serious, original and defensible model that is not disqualified on any of the grounds that usually disqualify wealth taxes, and that does not, at the rates tested, establish the full replacement of Capital Gains Tax and Inheritance Tax that this pass set as its own ceiling — provided it is presented with its conditions and limits intact, never oversold beyond them, and understood to await the government costing and counsel's opinion that only they can supply.

This pass used public data: HMRC receipts and OBR forecasts for the replacement ceiling; the Wealth Tax Commission and subsequent analyses for the base; and the European Court of Human Rights case-law, principally *Hentrich v France* (1994), for the durability tests. It is an order-of-magnitude and legal-principle assessment, not a costing or a legal opinion. A government costing requires taxpayer-level data this pass does not have; a definitive human-rights assessment requires counsel's opinion.

Changes in this version

Version 2, 29 August 2026. Two changes, one of them material.

The replacement ceiling has risen, and Test 1 no longer clears it. This pass adopted full replacement of Capital Gains Tax and Inheritance Tax as its revenue test. That was a deliberately severe condition, chosen because the burden of proof in an adversarial pass belongs on the model. It was not a requirement the architecture paper asserts: that paper presents replacement as a conditional possibility

and reserves revenue sufficiency as an open empirical question. HMRC's own statistics table puts CGT and IHT receipts for 2025-26 at £24.310bn and £8.467bn, a combined £32.8bn, against the £21–23bn used in version 1. On the same base and the same behavioural haircuts, the tested rate of 2 per cent no longer clears it, and the rate that would be in a region where the proportionality problem identified at Test 3 becomes materially worse. Test 1 and the verdicts resting on it, at Sections 3, 8, 9 and 10, have been revised accordingly.

One inference has been withdrawn. Version 1 read a fall in CGT receipts in 2024-25 as evidence of realisation-based lock-in. That fall was a payment-timing effect; liabilities for the same year rose sharply as holders made disposals ahead of anticipated policy changes. The observation has been reclassified as evidence of timing behaviour and no longer supports the lock-in claim. Section 1.1 also now states which statistical series this pass uses, because receipts and liabilities differ by £10.5bn for the same tax and year, and version 1 did not say which it meant.

Neither change alters the architecture paper's treatment of full replacement, which presents it as a conditional possibility rather than an asserted requirement. What has been corrected is this pass's own reporting, which allowed success against a self-imposed ceiling to be stated as a property of the model.

Separately, the architecture paper has been revised to Version 2 so that its Section 4.3 states the two Hentrich safeguards as requirements. That correction originates in Test 2 of this pass and is recorded in that paper.